



<https://hpc.gwdg.de>

Lars Quentin

MPI-based Creation and Benchmarking of a Dynamic Elasticsearch Cluster

1 Introduction

2 Spawner

3 Ingestor

4 Querier

5 Test Evaluation

6 Conclusion

Insights

- Why a custom spawner and new specialized benchmarker is required
- How the following works:
 - ▶ distributed cluster spawner
 - ▶ distributed ingestion benchmarker
 - ▶ distributed query benchmarker
- How to create a new benchmark scenario from scratch

Motivation: Data Lakes

Why are Data Lakes needed

- Research becomes evermore data-driven and compute-intensive
 - ▶ More Simulations
 - ▶ Data Science, Machine Learning
- HPC becomes more data oriented
- Better data-management tooling needed
- HPC operates on raw data
 - ⇒ Data Lakes

Motivation: Data Lakes

Why are Data Lakes needed

- Research becomes evermore data-driven and compute-intensive
 - ▶ More Simulations
 - ▶ Data Science, Machine Learning
- HPC becomes more data oriented
- Better data-management tooling needed
- HPC operates on raw data
 - ⇒ Data Lakes

Metadata management

- Providing storage is easy
- Managing storage is hard
- Keep data findable, manage data
- Fully indexed
- Fully (fuzzy) searchable
- No-SQL data store / search engine
 - ▶ Elasticsearch

Motivation: Elasticsearch and Rally

Elasticsearch for HPC

- Elasticsearch is designed for cloud-use
 - ▶ Always running
 - ▶ Same host, same IP
 - ▶ Only ethernet
- This is not given in HPC:
 - ▶ Jobs spawned on demand
 - ▶ Every job gets different nodes
 - ▶ Changing IPs between runs
 - ▶ ETH, IB, Intel OPA
- Thus, a custom **stateful** workflow is required for HPC use!

Motivation: Elasticsearch and Rally

Elasticsearch for HPC

- Elasticsearch is designed for cloud-use
 - ▶ Always running
 - ▶ Same host, same IP
 - ▶ Only ethernet
- This is not given in HPC:
 - ▶ Jobs spawned on demand
 - ▶ Every job gets different nodes
 - ▶ Changing IPs between runs
 - ▶ ETH, IB, Intel OPA
- Thus, a custom **stateful** workflow is required for HPC use!

Benchmarking Elasticsearch

- HPC is all about performance
- Elastic's benchmarker: *rally* [1]
 - ▶ Used for in-house performance regression testing
 - ▶ Written in Python
 - ▶ Distributed using *thespian* agent framework
 - ▶ After previous unpublished research at GWDG:
 - Doesn't work with over 60 nodes
- Not viable for HPC-scale benchmarking

Contributions

The 4 main contributions of this work:

Contributions

The 4 main contributions of this work:

1 On-demand Elasticsearch Cluster Spawner

- ▶ Zero-configuration
- ▶ Dynamic resolution, based on SLURM MPI environment
- ▶ Arbitrary cluster size
- ▶ Stateful between runs

Contributions

The 4 main contributions of this work:

1 On-demand Elasticsearch Cluster Spawner

- ▶ Zero-configuration
- ▶ Dynamic resolution, based on SLURM MPI environment
- ▶ Arbitrary cluster size
- ▶ Stateful between runs

2 Ingestion Benchmark

- ▶ Distributed, MPI-based
- ▶ Benchmarks easily portable from rally

Contributions

The 4 main contributions of this work:

1 On-demand Elasticsearch Cluster Spawner

- ▶ Zero-configuration
- ▶ Dynamic resolution, based on SLURM MPI environment
- ▶ Arbitrary cluster size
- ▶ Stateful between runs

2 Ingestion Benchmark

- ▶ Distributed, MPI-based
- ▶ Benchmarks easily portable from rally

3 Query Benchmark

- ▶ Distributed, MPI-based
- ▶ Mixed queries for realistic load
- ▶ Custom scenario support using own JSON-based DSL

Contributions

The 4 main contributions of this work:

1 On-demand Elasticsearch Cluster Spawner

- ▶ Zero-configuration
- ▶ Dynamic resolution, based on SLURM MPI environment
- ▶ Arbitrary cluster size
- ▶ Stateful between runs

2 Ingestion Benchmark

- ▶ Distributed, MPI-based
- ▶ Benchmarks easily portable from rally

3 Query Benchmark

- ▶ Distributed, MPI-based
- ▶ Mixed queries for realistic load
- ▶ Custom scenario support using own JSON-based DSL

4 Example workflow for canonical dataset

Background: Elasticsearch

- Distributed search engine
- Document-based NoSQL-Storage
- Internally based on Apache Lucene
- Provides JSON-based REST interface
- Apache 2.0 fork: Opensearch
- Advantages:
 - ▶ Mature ecosystem
 - ▶ Very battle-tested
 - ▶ A lot of tooling / library support



Background: Benchmarking

- For elasticsearch: All literature uses rally [2] [3] [4]
- Alternatives: Just use a HTTP benchmarker
 - ▶ JMeter [5]
 - ▶ wrk [6]
 - ▶ Grafana k6 [7]
- Most NoSQL comparisons are done by database vendors [8]
 - ▶ Bad financial incentives

On-Demand, Dynamic Cluster Spawner

Features

On-Demand, Dynamic Cluster Spawner

Features

- Fully automated, uses MPI environment provided by SLURM

On-Demand, Dynamic Cluster Spawner

Features

- Fully automated, uses MPI environment provided by SLURM
- Dynamically fetches the hosts
 - ▶ Not required to know them beforehand
 - ▶ IPs and hardware can be changed between runs

On-Demand, Dynamic Cluster Spawner

Features

- Fully automated, uses MPI environment provided by SLURM
- Dynamically fetches the hosts
 - ▶ Not required to know them beforehand
 - ▶ IPs and hardware can be changed between runs
- Very portable through containerization (Singularity)

On-Demand, Dynamic Cluster Spawner

Features

- Fully automated, uses MPI environment provided by SLURM
- Dynamically fetches the hosts
 - ▶ Not required to know them beforehand
 - ▶ IPs and hardware can be changed between runs
- Very portable through containerization (Singularity)
- Stateful: Same cluster can be respawned
 - ▶ on **different nodes**
 - ▶ **without reingestion**

On-Demand, Dynamic Cluster Spawner

Features

- Fully automated, uses MPI environment provided by SLURM
- Dynamically fetches the hosts
 - ▶ Not required to know them beforehand
 - ▶ IPs and hardware can be changed between runs
- Very portable through containerization (Singularity)
- Stateful: Same cluster can be respawned
 - ▶ on **different nodes**
 - ▶ **without reingestion**
- NIC-agnostic. Tested on:
 - ▶ Ethernet
 - ▶ Infiniband

On-Demand, Dynamic Cluster Spawner (cont.)

High-Level Workflow:

Prerequisites:

- All hosts are known to each other via the MPI environment
- All nodes have at least one shared mount

On-Demand, Dynamic Cluster Spawner (cont.)

High-Level Workflow:

Prerequisites:

- All hosts are known to each other via the MPI environment
- All nodes have at least one shared mount

Workflow:

- 1 Each node creates a config

On-Demand, Dynamic Cluster Spawner (cont.)

High-Level Workflow:

Prerequisites:

- All hosts are known to each other via the MPI environment
- All nodes have at least one shared mount

Workflow:

- 1 Each node creates a config
- 2 MPI-Gather all hostnames to the root rank

On-Demand, Dynamic Cluster Spawner (cont.)

High-Level Workflow:

Prerequisites:

- All hosts are known to each other via the MPI environment
- All nodes have at least one shared mount

Workflow:

- 1 Each node creates a config
- 2 MPI-Gather all hostnames to the root rank
- 3 The root node updates the configs for all nodes

On-Demand, Dynamic Cluster Spawner (cont.)

High-Level Workflow:

Prerequisites:

- All hosts are known to each other via the MPI environment
- All nodes have at least one shared mount

Workflow:

- 1 Each node creates a config
- 2 MPI-Gather all hostnames to the root rank
- 3 The root node updates the configs for all nodes
- 4 Each rank starts its singularity container with the config bind-mounted in

On-Demand, Dynamic Cluster Spawner (cont.)

High-Level Workflow:

Prerequisites:

- All hosts are known to each other via the MPI environment
- All nodes have at least one shared mount

Workflow:

- 1 Each node creates a config
- 2 MPI-Gather all hostnames to the root rank
- 3 The root node updates the configs for all nodes
- 4 Each rank starts its singularity container with the config bind-mounted in

See the accompanying report for a more low-level workflow.

On-Demand, Dynamic Cluster Spawner (cont.)

Example Generated Config

```
1  cluster.name: securemetadata
2  node.name: securemetadata4
3  node.roles: ["master", "data"]
4  network.host: 0.0.0.0
5  cluster.initial_master_nodes: [securemetadata0]
6  # Expects hostnames to be DNS resolvable
7  discovery.seed_hosts: [
8      "hostname_of_rank_0",
9      "hostname_of_rank_1",
10     "hostname_of_rank_2"
11  ]
12 xpack.security.enabled: false
```

Ingestion Benchmark

■ Two purposes:

- 1 Ingest JSON corpus into Elasticsearch cluster for query benchmarks
- 2 Measure performance of **write-performance** and throughput

■ Features:

- ▶ Distributed, MPI-based
- ▶ I/O optimized through *offset caching*
- ▶ Supports statically typed index definitions
- ▶ Supports Newline Delimited JSON (NDJSON)
 - Thus compatible with rally!
- ▶ Configurable via CLI: bulk size, shards per node

Offset Caching

Problem

- Data has to be partitioned.

Offset Caching

Problem

- Data has to be partitioned.
- This should be done fairly, i.e.

For N nodes and L lines, rank i gets

$$\left[\frac{i}{N} \cdot L, \frac{i+1}{N} \cdot L \right)$$

Offset Caching

Problem

- Data has to be partitioned.
- This should be done fairly, i.e.
For N nodes and L lines, rank i gets
 $\left[\frac{i}{N} \cdot L, \frac{i+1}{N} \cdot L \right)$
- For the computation the number of lines need to be known. (1st read)

Offset Caching

Problem

- Data has to be partitioned.
- This should be done fairly, i.e.
For N nodes and L lines, rank i gets
 $\left[\frac{i}{N} \cdot L, \frac{i+1}{N} \cdot L \right)$
- For the computation the number of lines need to be known. (1st read)
- Afterwards, the line has to be found. (2nd read)
 - ▶ Can't just seek, since JSON documents have variadic size!

Offset Caching

Problem

- Data has to be partitioned.
- This should be done fairly, i.e.
For N nodes and L lines, rank i gets
 $\left[\frac{i}{N} \cdot L, \frac{i+1}{N} \cdot L \right)$
- For the computation the number of lines need to be known. (1st read)
- Afterwards, the line has to be found. (2nd read)
 - ▶ Can't just seek, since JSON documents have variadic size!
- A lot of I/O, the corpus is 75GB.

Offset Caching

Problem

- Data has to be partitioned.
- This should be done fairly, i.e.
For N nodes and L lines, rank i gets
 $\left[\frac{i}{N} \cdot L, \frac{i+1}{N} \cdot L \right)$
- For the computation the number of lines need to be known. (1st read)
- Afterwards, the line has to be found. (2nd read)
 - ▶ Can't just seek, since JSON documents have variadic size!
- A lot of I/O, the corpus is 75GB.

Solution

- Just one node computes it, and caches it in a file!

Offset Caching

Problem

- Data has to be partitioned.
- This should be done fairly, i.e.
For N nodes and L lines, rank i gets $\left[\frac{i}{N} \cdot L, \frac{i+1}{N} \cdot L\right)$
- For the computation the number of lines need to be known. (1st read)
- Afterwards, the line has to be found. (2nd read)
 - ▶ Can't just seek, since JSON documents have variadic size!
- A lot of I/O, the corpus is 75GB.

Solution

- Just one node computes it, and caches it in a file!
- Steps:
 - 1 Read 1: Count number of lines.
 - 2 Compute starting and ending line for each rank.
 - 3 Read 2: Find the byte offsets for each rank.
 - 4 Save everything into a `.offsets.json` file.

Offset Caching (cont.)

Example .offset.json file for 3 nodes

```
1  {
2    "number_of_workers":3,
3    "offsets":[
4      {
5        "rank":0,
6        "starting_line":0,
7        "starting_byte":0,
8        "number_of_lines":8333
9      },
10     { "rank":1, "starting_line":8333,
11       "starting_byte":4157901, "number_of_lines":8333 },
12     { "rank":2, "starting_line":16666,
13       "starting_byte":8315734, "number_of_lines":null }
14   ] }
```

Ingestion Benchmark (cont.)

Workflow: Setup (root only)

- Create cache offsets if not already existing
 - ▶ Requires same number of load generators

Ingestion Benchmark (cont.)

Workflow: Setup (root only)

- Create cache offsets if not already existing
 - ▶ Requires same number of load generators
- Create empty Elasticsearch index with following settings:
 - ▶ Strict type mappings (Elasticsearch syntax)
 - ▶ One shard per Cluster node (configurable)
 - ▶ `requests.cache.enable:`
`false`

Ingestion Benchmark (cont.)

Workflow: Setup (root only)

- Create cache offsets if not already existing
 - ▶ Requires same number of load generators
- Create empty Elasticsearch index with following settings:
 - ▶ Strict type mappings (Elasticsearch syntax)
 - ▶ One shard per Cluster node (configurable)
 - ▶ `requests.cache.enable:`
`false`

Workflow: Benchmark

- Each rank chooses one ES node to send to

Ingestion Benchmark (cont.)

Workflow: Setup (root only)

- Create cache offsets if not already existing
 - ▶ Requires same number of load generators
- Create empty Elasticsearch index with following settings:
 - ▶ Strict type mappings (Elasticsearch syntax)
 - ▶ One shard per Cluster node (configurable)
 - ▶ `requests.cache.enable: false`

Workflow: Benchmark

- Each rank chooses one ES node to send to
- Seek to starting byte based on offsets

Ingestion Benchmark (cont.)

Workflow: Setup (root only)

- Create cache offsets if not already existing
 - ▶ Requires same number of load generators
- Create empty Elasticsearch index with following settings:
 - ▶ Strict type mappings (Elasticsearch syntax)
 - ▶ One shard per Cluster node (configurable)
 - ▶ `requests.cache.enable: false`

Workflow: Benchmark

- Each rank chooses one ES node to send to
- Seek to starting byte based on offsets
- Send the requests blockingly, as fast as possible

Ingestion Benchmark (cont.)

Workflow: Setup (root only)

- Create cache offsets if not already existing
 - ▶ Requires same number of load generators
- Create empty Elasticsearch index with following settings:
 - ▶ Strict type mappings (Elasticsearch syntax)
 - ▶ One shard per Cluster node (configurable)
 - ▶ `requests.cache.enable: false`

Workflow: Benchmark

- Each rank chooses one ES node to send to
- Seek to starting byte based on offsets
- Send the requests blockingly, as fast as possible
- Track response time *directly* after

Ingestion Benchmark (cont.)

Workflow: Setup (root only)

- Create cache offsets if not already existing
 - ▶ Requires same number of load generators
- Create empty Elasticsearch index with following settings:
 - ▶ Strict type mappings (Elasticsearch syntax)
 - ▶ One shard per Cluster node (configurable)
 - ▶ `requests.cache.enable: false`

Workflow: Benchmark

- Each rank chooses one ES node to send to
- Seek to starting byte based on offsets
- Send the requests blockingly, as fast as possible
- Track response time *directly* after
- Wait at barrier

Ingestion Benchmark (cont.)

Workflow: Setup (root only)

- Create cache offsets if not already existing
 - ▶ Requires same number of load generators
- Create empty Elasticsearch index with following settings:
 - ▶ Strict type mappings (Elasticsearch syntax)
 - ▶ One shard per Cluster node (configurable)
 - ▶ `requests.cache.enable: false`

Workflow: Benchmark

- Each rank chooses one ES node to send to
- Seek to starting byte based on offsets
- Send the requests blockingly, as fast as possible
- Track response time *directly* after
- Wait at barrier
- MPI Gather all data at root, dump into JSON file

Query Benchmark

- Measures query-/**read**-performance against previously ingested data.

Query Benchmark

- Measures query-/**read**-performance against previously ingested data.
- Works through scenarios in a fork-join model.
 - ▶ Supports mixing queries in same scenario

Query Benchmark

- Measures query-/**read**-performance against previously ingested data.
- Works through scenarios in a fork-join model.
 - ▶ Supports mixing queries in same scenario
- Features:
 - ▶ Distributed, MPI-based

Query Benchmark

- Measures query-/**read**-performance against previously ingested data.
- Works through scenarios in a fork-join model.
 - ▶ Supports mixing queries in same scenario
- Features:
 - ▶ Distributed, MPI-based
 - ▶ Fully configurable by JSON-DSL; no hard-coded scenarios
 - No need to edit the source code
 - Embeds Elasticsearch syntax internally ⇒ accessible for ES-users
 - Simplification of Rally syntax ⇒ easy to port

Query Benchmark

- Measures query-/**read**-performance against previously ingested data.
- Works through scenarios in a fork-join model.
 - ▶ Supports mixing queries in same scenario
- Features:
 - ▶ Distributed, MPI-based
 - ▶ Fully configurable by JSON-DSL; no hard-coded scenarios
 - No need to edit the source code
 - Embeds Elasticsearch syntax internally ⇒ accessible for ES-users
 - Simplification of Rally syntax ⇒ easy to port
 - ▶ Bypasses the cache

Query Benchmark

- Measures query-/**read**-performance against previously ingested data.
- Works through scenarios in a fork-join model.
 - ▶ Supports mixing queries in same scenario
- Features:
 - ▶ Distributed, MPI-based
 - ▶ Fully configurable by JSON-DSL; no hard-coded scenarios
 - No need to edit the source code
 - Embeds Elasticsearch syntax internally ⇒ accessible for ES-users
 - Simplification of Rally syntax ⇒ easy to port
 - ▶ Bypasses the cache
 - ▶ Parses the responses for more data; not only based on HTTP response code

Query Benchmark

- Measures query-/**read**-performance against previously ingested data.
- Works through scenarios in a fork-join model.
 - ▶ Supports mixing queries in same scenario
- Features:
 - ▶ Distributed, MPI-based
 - ▶ Fully configurable by JSON-DSL; no hard-coded scenarios
 - No need to edit the source code
 - Embeds Elasticsearch syntax internally ⇒ accessible for ES-users
 - Simplification of Rally syntax ⇒ easy to port
 - ▶ Bypasses the cache
 - ▶ Parses the responses for more data; not only based on HTTP response code
 - ▶ Test mode for easier debugging

Input Format for Query Benchmark (part 1)

```
1  [  
2    {  
3      "search_queries": [  
4        {  
5          /* everything in here just gets sent to ES */  
6          "body": {  
7            /* The raw ES query sent to the server */  
8          }  
9        }  
10     ],  
11     "warmup_time_secs": 30, /* optional */  
12     "execution_time_secs": 120, /* optional */  
13   },  
14   ...
```

Input Format for Query Benchmark (part 2)

```
1  {
2    "search_queries": [
3      {
4        "body": {
5          /* The first of 2 queries sent iteratively (random order) */
6        }
7      },
8      {
9        "body": {
10         /* The second of 2 queries sent iteratively (random order) */
11       }
12     }
13   ],
14   "warmup_time_secs": 30, /* optional */
15   "execution_time_secs": 180, /* optional */
16   "sleep_between_requests_secs": 0.25 /* optional */
17 }
```

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green
- Partition the hosts onto the load generators

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green
- Partition the hosts onto the load generators
- If warmup time is set: Send queries, discard result (fill OS caches)

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green
- Partition the hosts onto the load generators
- If warmup time is set: Send queries, discard result (fill OS caches)
- After that, until execution time for current step is reached:
 - ▶ Select next query in current step

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green
- Partition the hosts onto the load generators
- If warmup time is set: Send queries, discard result (fill OS caches)
- After that, until execution time for current step is reached:
 - ▶ Select next query in current step
 - ▶ Track before, send query, track after

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green
- Partition the hosts onto the load generators
- If warmup time is set: Send queries, discard result (fill OS caches)
- After that, until execution time for current step is reached:
 - ▶ Select next query in current step
 - ▶ Track before, send query, track after
 - ▶ Parse ES response, save (latency, docs count)

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green
- Partition the hosts onto the load generators
- If warmup time is set: Send queries, discard result (fill OS caches)
- After that, until execution time for current step is reached:
 - ▶ Select next query in current step
 - ▶ Track before, send query, track after
 - ▶ Parse ES response, save (latency, docs count)
 - ▶ Sleep if configured

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green
- Partition the hosts onto the load generators
- If warmup time is set: Send queries, discard result (fill OS caches)
- After that, until execution time for current step is reached:
 - ▶ Select next query in current step
 - ▶ Track before, send query, track after
 - ▶ Parse ES response, save (latency, docs count)
 - ▶ Sleep if configured
- Wait at MPI barrier for next step

Query Benchmark (cont.)

High-Level Workflow

For each disjunct fork-join benchmark step:

- Wait for Elasticsearch cluster health to be green
- Partition the hosts onto the load generators
- If warmup time is set: Send queries, discard result (fill OS caches)
- After that, until execution time for current step is reached:
 - ▶ Select next query in current step
 - ▶ Track before, send query, track after
 - ▶ Parse ES response, save (latency, docs count)
 - ▶ Sleep if configured
- Wait at MPI barrier for next step

See the accompanying report for a more low-level workflow.

How to Create and Run a Benchmark

- 1 Choose a dataset or create a synthetic one
 - ▶ format as NDJSON

How to Create and Run a Benchmark

- 1 Choose a dataset or create a synthetic one
 - ▶ format as NDJSON
- 2 Define the Elasticsearch type mappings for each attribute

How to Create and Run a Benchmark

- 1 Choose a dataset or create a synthetic one
 - ▶ format as NDJSON
- 2 Define the Elasticsearch type mappings for each attribute
- 3 Design the query document
 - ▶ Basically just embedding the Elasticsearch API queries into more JSON
 - ▶ Note: They can thus be easily tested using cURL/Postman/Insomnia/...

How to Create and Run a Benchmark

- 1 Choose a dataset or create a synthetic one
 - ▶ format as NDJSON
- 2 Define the Elasticsearch type mappings for each attribute
- 3 Design the query document
 - ▶ Basically just embedding the Elasticsearch API queries into more JSON
 - ▶ Note: They can thus be easily tested using cURL/Postman/Insomnia/...
- 4 Spawn up the cluster using SLURMs MPI environment

How to Create and Run a Benchmark

- 1 Choose a dataset or create a synthetic one
 - ▶ format as NDJSON
- 2 Define the Elasticsearch type mappings for each attribute
- 3 Design the query document
 - ▶ Basically just embedding the Elasticsearch API queries into more JSON
 - ▶ Note: They can thus be easily tested using cURL/Postman/Insomnia/...
- 4 Spawn up the cluster using SLURMs MPI environment
- 5 Run the distributed ingestor to ingest the NDJSON corpus

How to Create and Run a Benchmark

- 1 Choose a dataset or create a synthetic one
 - ▶ format as NDJSON
- 2 Define the Elasticsearch type mappings for each attribute
- 3 Design the query document
 - ▶ Basically just embedding the Elasticsearch API queries into more JSON
 - ▶ Note: They can thus be easily tested using cURL/Postman/Insomnia/...
- 4 Spawn up the cluster using SLURMs MPI environment
- 5 Run the distributed ingestor to ingest the NDJSON corpus
- 6 Run the distributed query benchmarker using the query document

How to Create and Run a Benchmark

- 1 Choose a dataset or create a synthetic one
 - ▶ format as NDJSON
- 2 Define the Elasticsearch type mappings for each attribute
- 3 Design the query document
 - ▶ Basically just embedding the Elasticsearch API queries into more JSON
 - ▶ Note: They can thus be easily tested using cURL/Postman/Insomnia/...
- 4 Spawn up the cluster using SLURMs MPI environment
- 5 Run the distributed ingestor to ingest the NDJSON corpus
- 6 Run the distributed query benchmarker using the query document
- 7 Analyze the output JSON using a language of your choice
Python example can be found in the Git repo.

Benchmark

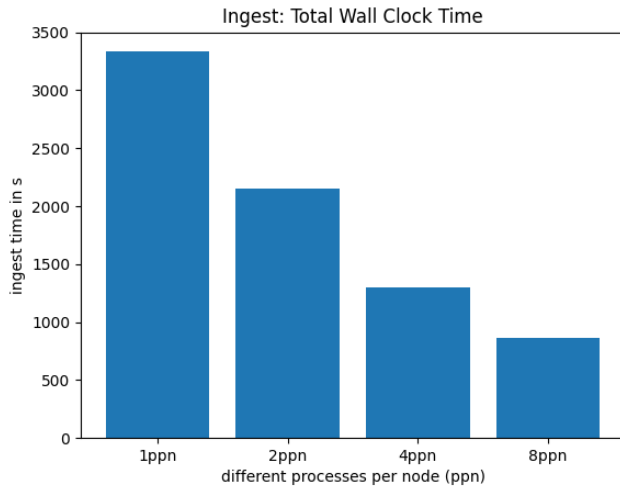
Dataset: NYC Taxis [9]

- All yellow taxi rides in NYC in 2015
- Published by NYC Taxi and Limousine Commission [10]
- 165 million documents, over 75GB
- Also used by Rally (Elastic)
- Most used for scaling testing
- Big documents, but mostly numeric data.


Setup

- 3 nodes on Emmy
- Ethernet
- Ubuntu 22.04 dockerhub image in Singularity
- Elasticsearch 8.11.0 with OpenJDK 21.0.1
- Python 3.9
- OpenMPI 4.1


Results




Results



Results



Results



Challenges/Open Problems

- Limited response size, hard limit by Elasticsearch's architecture
- Not possible to map load generator to cluster node according to optimal network topology
- Load generators and clusters cant share the same node
- Elasticsearch requires a custom kernel setting

Summary

- Project was a success, fully implemented both workflow and benchmarker
- Zero configuration needed once the benchmark was initially designed
- Fully integrated into SLURM
- Contributions:
 - 1 On-demand Elasticsearch Cluster Spawner
 - 2 Ingestion Benchmark
 - 3 Query Benchmark
 - 4 Example workflow for canonical dataset

References I

Elasticsearch Benchmarks. URL: <https://elasticsearch-benchmarks.elastic.co/#> (visited on 01/17/2024).

Hårek Haugerud, Mohamad Sobhie, and Anis Yazidi. “Tuning of Elasticsearch Configuration: Parameter Optimization Through Simultaneous Perturbation Stochastic Approximation”. In: *Frontiers in Big Data* 5 (2022). ISSN: 2624-909X. URL: <https://www.frontiersin.org/articles/10.3389/fdata.2022.686416> (visited on 01/18/2024).

Wataru Takase et al. *A solution for secure use of Kibana and Elasticsearch in multi-user environment*. June 30, 2017. arXiv: 1706.10040[cs]. URL: <http://arxiv.org/abs/1706.10040> (visited on 01/18/2024).

Hui Dou, Pengfei Chen, and Zibin Zheng. “Hdconfigor: Automatically Tuning High Dimensional Configuration Parameters for Log Search Engines”. In: *IEEE Access* 8 (2020), pp. 80638–80653. ISSN: 2169-3536. DOI: 10.1109/ACCESS.2020.2990735. URL: <https://ieeexplore.ieee.org/document/9079492/> (visited on 01/18/2024).

Apache JMeter - Apache JMeter™. URL: <https://jmeter.apache.org/> (visited on 01/18/2024).

Will Glozer. *wg/wrk*. original-date: 2012-03-20T11:12:28Z. Jan. 18, 2024. URL: <https://github.com/wg/wrk> (visited on 01/18/2024).

References II

Load testing for engineering teams | Grafana k6. URL: <https://k6.io> (visited on 01/18/2024).

timescale/tsbs. original-date: 2018-08-08T14:30:28Z. Jan. 17, 2024. URL:
<https://github.com/timescale/tsbs> (visited on 01/18/2024).

rally-tracks/nyc_taxis at master · elastic/rally-tracks. GitHub. URL:
https://github.com/elastic/rally-tracks/tree/master/nyc_taxis (visited on 02/15/2024).

TLC Trip Record Data - TLC. URL:
<https://www.nyc.gov/site/tlc/about/tlc-trip-record-data.page> (visited on 02/15/2024).